



SAPIENZA  
UNIVERSITÀ DI ROMA

# Input-Sensitive Profiling

Emilio Coppa

June 1, 2012

Ingegneria degli Algoritmi


Corso di Laurea in Informatica – A.A. 2011-2012

# Outline

- 1 Introduction
  - Conventional profilers
  - Drawbacks classical approach
- 2 Read Memory Size
  - Our approach
  - Definition
  - Examples
- 3 Case study: wf
- 4 Algorithm
- 5 Implementation
- 6 Experiments: SPEC CPU2006

# Conventional profilers


Conventional profilers gather **cumulative** info over a whole execution



⇒ No information about how single portions of the code **scale** as a function of **input size**

# Drawbacks classical approach

- Often hard to extract portions of code from an application and analyze them separately
- Hard to collect real data about typical usage scenarios to be reproduced in experiments
- Miss cache effects due to interaction with the overall application



Critical algorithmic code should be analyzed within the **actual** context of applications it is deployed in

# Our approach

**“Input-Sensitive”** profiling: aggregating routine times by input sizes

For each routine  $f$ , collect a tuple:

$$\langle n_i, c_i, \max_i, \min_i, \text{sum}_i, q_i \rangle$$

for each distinct value of the input size, where:

- $n_i$  = estimate of an input size
- $c_i$  = # of times the routine is called on input size  $n_i$
- $\max_i / \min_i$  = maximum and minimum costs required by any execution of  $f$  on input size  $n_i$
- $\text{sum}_i / q_i$  = sum of the costs required by the executions of  $f$  on input size  $n_i$  and the sum of the costs' squares

# How to measure input size automatically?

## Input size $\approx$ Read Memory Size

The **read memory size** (RMS) of the execution of a routine  $f$  is the number of distinct memory cells first accessed by  $f$ , or by a descendant of  $f$  in the call tree, with a read operation.

# Read Memory Size (Example)

```
void swap(int * a, int * b) {  
    int temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

The function `swap` has RMS 2 because it reads (first access) objects `*a` and `*b`, and writes (first access) variable `temp`

# Read Memory Size (Example 2)

```
call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
  return
read w
return
```

Fn	Accessed cells (first-read green)	RMS



# Read Memory Size (Example 2)

```

→ call f
    read x
    write y
    call g
        read x
        read y
        read z
        write w
    return
read w
return

```

Fn	Accessed cells (first-read green)	RMS
f		

# Read Memory Size (Example 2)

```

→ call f
    read x
    write y
    call g
        read x
        read y
        read z
        write w
    return
read w
return
  
```

Fn	Accessed cells (first-read green)	RMS
f	x	

# Read Memory Size (Example 2)

```

→ call f
    read x
    write y
    call g
        read x
        read y
        read z
        write w
    return
    read w
    return
  
```

Fn	Accessed cells (first-read green)	RMS
f	<span style="background-color: #90EE90;">x</span> y	

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
  return
read w
return
  
```

→

Fn	Accessed cells (first-read green)	RMS
f	<span style="background-color: #90EE90;">x</span> y	
g		

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    →   read x
        read y
        read z
        write w
      return
  read w
  return

```

Fn	Accessed cells (first-read green)	RMS
f	<span style="border: 1px solid green; border-radius: 10px; padding: 2px;">x</span> y	
g	<span style="border: 1px solid green; border-radius: 10px; padding: 2px;">x</span>	

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
  return
read w
return

```



Fn	Accessed cells (first-read green)	RMS
f	<span style="border: 1px solid green; border-radius: 5px; padding: 2px;">x</span> y	
g	<span style="border: 1px solid green; border-radius: 5px; padding: 2px;">x</span> <span style="border: 1px solid green; border-radius: 5px; padding: 2px;">y</span>	

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
  return
read w
return

```



Fn	Accessed cells (first-read green)	RMS
f	<span style="border: 1px solid green; border-radius: 5px; padding: 2px;">x</span> y <span style="border: 1px solid green; border-radius: 5px; padding: 2px;">z</span>	
g	<span style="border: 1px solid green; border-radius: 5px; padding: 2px;">x</span> <span style="border: 1px solid green; border-radius: 5px; padding: 2px;">y</span> <span style="border: 1px solid green; border-radius: 5px; padding: 2px;">z</span>	

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
    return
  read w
return

```



Fn	Accessed cells (first-read green)	RMS
f	<span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">x</span> y <span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">z</span> w	
g	<span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">x</span> <span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">y</span> <span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">z</span> w	



# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
    return
  read w
  return

```



F <sub>n</sub>	Accessed cells (first-read green)	RMS
f	<span style="background-color: #90EE90;">x</span> y <span style="background-color: #90EE90;">z</span> w	3
g	<span style="background-color: #90EE90;">x</span> <span style="background-color: #90EE90;">y</span> <span style="background-color: #90EE90;">z</span> w	

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
    return
  → read w
  return

```

F <sub>n</sub>	Accessed cells (first-read green)	RMS
f	<span style="background-color: #90EE90;">x</span> y <span style="background-color: #90EE90;">z</span> w	3
g	<span style="background-color: #90EE90;">x</span> <span style="background-color: #90EE90;">y</span> <span style="background-color: #90EE90;">z</span> w	

# Read Memory Size (Example 2)

```

call f
  read x
  write y
  call g
    read x
    read y
    read z
    write w
    return
  read w
→ return

```

Fn	Accessed cells (first-read green)	RMS
f	<span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">x</span> y <span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">z</span> w	2
g	<span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">x</span> <span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">y</span> <span style="background-color: #d9ead3; border: 1px solid black; padding: 2px;">z</span> w	3

# Case study: wf


We discuss `wf-0.41`, a simple word frequency counter included in the current development head of Linux Fedora (Fedora 17–Beefy Miracle).

We profile `wf` with:



`gprof` a traditional and well-known call graph execution profiler – <http://www.gnu.org/software/binutils/>

`aprof` our implementation of an input-sensitive profiler – <http://code.google.com/p/aprof/>

## gprof vs aprof



## gprof vs aprof



# Is there any bottleneck in str\_tolower?

```
void str_tolower(char* str) {  
    int i;  
    for (i = 0; i < strlen(str); i++)  
        str[i] = wf_tolower(str[i]);  
}
```

# Is there any bottleneck in str\_tolower?

```
void str_tolower(char* str) {  
    int i;  
    for (i = 0; i < strlen(str); i++)  
        str[i] = wf_tolower(str[i]);  
}
```



# Is there any bottleneck in str\_tolower?

```
void str_tolower(char* str) {  
    int i;  
    for (i = 0; i < strlen(str); i++)  
        str[i] = wf_tolower(str[i]);  
}
```

Why did gprof **fail** to reveal the quadratic trend of str\_tolower?

# Short vs long words

Input: Anna Karenina

52.2%	<code>addword</code>
31.3%	<code>str_tolower</code>



Input: Protein sequences

61.8%	<code>str_tolower</code>
32.6%	<code>addword</code>

- Need to have different workloads for different routines!
- How do we know in advance which routine is a bottleneck?


# Fixing the code

Loop invariant code motion:

```
void str_tolower(char* str) {  
    int i;  
    int len = strlen(str);  
    for (i = 0; i < len; i++)  
        str[i] = wf_tolower(str[i]);  
}
```

Improvements:	6%	Anna Karenina
	30%	Protein sequences

# Computing RMS (data structures)



Shadow run-time  
stack  $S$

For each  $i \in [0, top]$ , the  $i$ -th stack entry  $S[i]$  stores:

- $rtn$ : id of the routine
- $ts$ : timestamp assigned to this activation
- $cost$ : cumulative cost
- $rms$ : *partial read memory size* of the activation

Each memory location  $w$  has a timestamp  $ts[w]$  which contains the time of the **latest** access to  $w$

# Computing RMS (algorithm)

**procedure** call( $r$ ):

$top++$

$S[top].rtn \leftarrow r$

$S[top].ts \leftarrow count$

$S[top].rms \leftarrow 0$

$S[top].cost \leftarrow get\_cost()$

**procedure** return():

$collect(S[top].rtn, S[top].rms, get\_cost() - S[top].cost)$

$S[top-1].rms += S[top].rms$

$top--$

# Computing RMS (algorithm)

```

procedure read( $w$ ):
  if  $ts[w] < S[top].ts$  then
     $S[top].rms++$ 
    if  $ts[w] \neq 0$  then
      let  $i$  be the max index in  $S$ 
      such that  $S[i].ts \leq ts[w]$ 
       $S[i].rms--$ 
    end if
  end if
   $ts[w] \leftarrow count$ 

procedure write( $w$ ):
   $ts[w] \leftarrow count$ 

```

# Computing RMS (example)

```
call f
```

```
    read x
    write y
    write z
    call g
        read x
        read y
        read z
        return
    return
```

**Step 1:**

```
S[0].id = f
S[0].cost = 0
S[0].ts = 1
S[0].rms = 0
```

```
ts[x] = 0
ts[y] = 0
ts[z] = 0
```

# Computing RMS (example)

```
call f
  read x
  write y
  write z
  call g
    read x
    read y
    read z
    return
  return
```

## Step 2:

```
S[0].id = f
S[0].cost = 0
S[0].ts = 1
S[0].rms = 1

ts[x] = 1
ts[y] = 0
ts[z] = 0
```



# Computing RMS (example)

```
call f
  read x
  write y
  write z
  call g
    read x
    read y
    read z
    return
  return
```

## Step 3:

```
S[0].id = f
S[0].cost = 0
S[0].ts = 1
S[0].rms = 1
```

```
ts[x] = 1
```

```
ts[y] = 1
```

```
ts[z] = 1
```

# Computing RMS (example)

```

call f
  read x
  write y
  write z
  call g
    read x
    read y
    read z
    return
  return

```

## Step 4:

```

S[1].id = g
S[1].cost = 4
S[1].ts = 2
S[1].rms = 0

```

```

S[0].id = f
S[0].cost = 0
S[0].ts = 1
S[0].rms = 1

```

```

ts[x] = 1
ts[y] = 1
ts[z] = 1

```

# Computing RMS (example)

```

call f
  read x
  write y
  write z
  call g
    read x
    read y
    read z
    return
  return

```

## Step 5:

```

S[1].id = g
S[1].cost = 4
S[1].ts = 2
S[1].rms = 3

S[0].id = f
S[0].cost = 0
S[0].ts = 1
S[0].rms = -2

```

```

ts[x] = 2
ts[y] = 2
ts[z] = 2

```

# Computing RMS (example)

```
call f
  read x
  write y
  write z
  call g
    read x
    read y
    read z
    return
  return
```

## Step 6:

$S[0].id = f$

$S[0].cost = 0$

$S[0].ts = 1$

$S[0].rms = -2 + 3 = 1$

$ts[x] = 2$

$ts[y] = 2$

$ts[z] = 2$

# Implementation




A dynamic instrumentation infrastructure that translates the binary code into an architecture-neutral intermediate representation (VEX)

Events	Instrumentation	Data structures
memory accesses	easy	shadow memory
threads	easy	thread state
function calls/returns	hard	shadow stack

# Implementation: shadow memory

To maintain the timestamps  $ts$  of memory cells needed for computing the RMS values, we shadow each memory location accessed by the program with a 32-bit counter and we use a two-levels lookup table:



# Implementation: memory tracing resolution

To reduce the space needed by the lookup table, aprof allows it to configure the resolution  $k$  of distinct observable memory objects:

- $k = 1$ : finest resolution
- $k = 2$ : 2-bytes words  $\rightarrow$  half of timestamps
- $k = 4$ : default, 4-bytes words  $\rightarrow$  1/4 of timestamps!
- ...

$\Rightarrow$  this can impact the # of distinct RMS observed!

## SPEC CPU2006 – Time (slowdown)

	memcheck	callgrind-base	callgrind-cache	aprof
CINT	15.7×	46.5×	98.8×	31.8×
CFP	21.3×	20.4×	92.7×	27.9×

memcheck does not trace function calls/returns

callgrind-base does not trace memory accesses

⇒ aprof delivers comparable performance wrt other Valgrind tools

⇒ single run with aprof  $\approx$  N runs with gprof!



## SPEC CPU2006 – Space (overhead)



	memcheck	callgrind-base	callgrind-cache	aprof
CINT	1.8×	1.3×	1.3×	2.2×
CFP	1.5×	1.3×	1.3×	1.9×

callgrind-base does not use a shadow memory

memcheck applies different compression schemes

# Experimental evaluation

How many performance tuples can be automatically collected for each routine from a *single run* of a program on a typical workload?



## Demo

Thanks!  
Questions?

